

Multi-label Streams, Concept Drift, and Sequential Data

Jesse Read



[D&K] IoT Stream Data Mining 2017-2018

December 20, 2017


Outline

- 1 Multi-labelled Streaming Data
- 2 k NN for Streams
- 3 Batch-incremental Ensemble
- 4 Incremental base learner
- 5 Hoeffding Trees
- 6 Neural Networks
- 7 Concept Drift
- 8 Temporal Dependence
- 9 Connections to Sequential Data
- 10 Unlabelled Instances
- 11 Summary

Multi-label Data Streams

Many applications, e.g.,

- text (email, twitter, web, social networks, ...)
- images (and video)
- audio
- sensory data (IoT, ...)
- reinforcement learning (agent in an online environment)
- ...



Implications of Data Streams

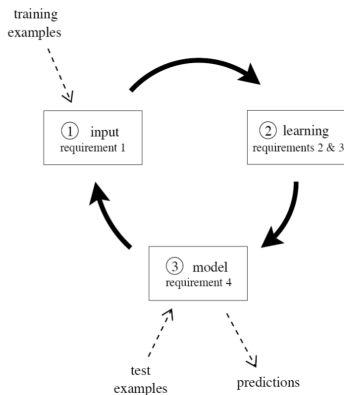
- 1 A **potentially infinite** sequence, $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_\infty$, but **limited memory**
- 2 Prediction \hat{y}_t must be made **immediately** at time t
- 3 Average time spent on prediction and learning per-instance – must be less than the average real intra-instance time (i.e., the real time between some \mathbf{x}_t and \mathbf{x}_{t+1})
- 4 Expect **concept drift** to occur: the distribution $(\mathbf{x}_t, y_t) \sim p_t(X, Y)$ changes with t .



Data-Stream Classification

Typical loop for data-stream learning, classification, and evaluation. For each instance:


- 1 Observe \mathbf{x}
- 2 Make a prediction $\hat{y} = h(\mathbf{x})$
- 3 Observe true label y
- 4 Measure the error $\epsilon = E(y, \hat{y})$
- 5 Monitor ϵ signal for concept drift
- 6 Update the model h with example (\mathbf{x}, y)



Implications for Multi-label Classifiers

As compared to the single-labelled case:


- More computationally complex (label dimension)
- Dynamic label set
- Dynamic label dependence
(i.e., multi-dimensional concept drift)
- Multi-labelled examples more difficult to obtain
(L -times more expensive)



Multi-label Learning in Data Streams

A good recipe to perform better than baseline binary relevance method in multi-label classification:


- 1 Analyze label dependence
- 2 Search for a good structure
- 3 Deploy model



Multi-label Learning in Data Streams

A good recipe to perform better than baseline binary relevance method in multi-label classification:

- 1 Analyze label dependence
- 2 Search for a good structure
- 3 Deploy model



No time! Data still arriving! Dependence (and therefore best structure) may become invalidated over time. Model should have been deployed already!


Multi-label Streams Methods

- k -Nearest Neighbours
- A **batch-incremental** ensemble
- Problem transformation with an **incremental base learner**
- Decision Rules
- Hoeffding trees
- Neural networks

Outline

- 1 Multi-labelled Streaming Data
- 2 *k*NN for Streams
- 3 Batch-incremental Ensemble
- 4 Incremental base learner
- 5 Hoeffding Trees
- 6 Neural Networks
- 7 Concept Drift
- 8 Temporal Dependence
- 9 Connections to Sequential Data
- 10 Unlabelled Instances
- 11 Summary

Incremental (Multi-label) k NN



k NN is a lazy method, just need to store a batch of instances;

- Compare new example to its neighbours in the batch, make a classification (see slides from last time)
- Add new instances (when labels are available) to the batch, purge out instances (e.g., FIFO) when search is too slow or memory is full.

Outline

- 1 Multi-labelled Streaming Data
- 2 k NN for Streams
- 3 Batch-incremental Ensemble**
- 4 Incremental base learner
- 5 Hoeffding Trees
- 6 Neural Networks
- 7 Concept Drift
- 8 Temporal Dependence
- 9 Connections to Sequential Data
- 10 Unlabelled Instances
- 11 Summary

Batch-Incremental Ensemble

Build regular models on batches/windows of instances (typically in a [weighted] ensemble).

$$\underbrace{\begin{bmatrix} \mathbf{x} \\ \mathbf{y} \end{bmatrix}_1, \begin{bmatrix} \mathbf{x} \\ \mathbf{y} \end{bmatrix}_2, \begin{bmatrix} \mathbf{x} \\ \mathbf{y} \end{bmatrix}_3, \begin{bmatrix} \mathbf{x} \\ \mathbf{y} \end{bmatrix}_4}_{\text{build } \mathbf{h}_1}, \underbrace{\begin{bmatrix} \mathbf{x} \\ \mathbf{y} \end{bmatrix}_5, \begin{bmatrix} \mathbf{x} \\ \mathbf{y} \end{bmatrix}_6, \begin{bmatrix} \mathbf{x} \\ \mathbf{y} \end{bmatrix}_7, \begin{bmatrix} \mathbf{x} \\ \mathbf{y} \end{bmatrix}_8}_{\text{build } \mathbf{h}_2}, \begin{bmatrix} \mathbf{x} \\ \mathbf{y} \end{bmatrix}_9, \begin{bmatrix} \mathbf{x} \\ \hat{\mathbf{y}} \end{bmatrix}_{10}$$

Using a batch size of w examples:


- build model \mathbf{h}_1 on $\{(\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_w, \mathbf{y}_w)\}$
- build model \mathbf{h}_2 on $\{(\mathbf{x}_{w+1}, \mathbf{y}_{w+1}), \dots, (\mathbf{x}_{2w}, \mathbf{y}_{2w})\}$
- build model \mathbf{h}_3 on $\{(\mathbf{x}_{2w+1}, \mathbf{y}_{2w+1}), \dots, (\mathbf{x}_{3w}, \mathbf{y}_{3w})\}$
- ...

Use the most recent M models; **predict** via vote :

$$\hat{\mathbf{y}} = \sum_{m=1}^M \mathbf{h}_m(\mathbf{x})$$

Batch-Incremental Ensemble

Build regular models on batches/windows of instances (typically in a [weighted] ensemble).



Using a batch size of w examples:


- build model \mathbf{h}_1 on $\{(\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_w, \mathbf{y}_w)\}$
- build model \mathbf{h}_2 on $\{(\mathbf{x}_{w+1}, \mathbf{y}_{w+1}), \dots, (\mathbf{x}_{2w}, \mathbf{y}_{2w})\}$
- build model \mathbf{h}_3 on $\{(\mathbf{x}_{2w+1}, \mathbf{y}_{2w+1}), \dots, (\mathbf{x}_{3w}, \mathbf{y}_{3w})\}$
- ...

Use the most recent M models; **predict** via **weighted** vote (more recent = more relevant):

$$\hat{\mathbf{y}} = \sum_{m=1}^M \omega_m \cdot \mathbf{h}_m(\mathbf{x}) \quad \triangleright \text{where } \omega_m \propto m$$

Batch-Incremental Ensemble


Build regular models on batches/windows of instances (typically in a [weighted] ensemble).



- A common approach, surprisingly effective
- Open choice of base classifier (e.g., C4.5, SVM, ...)
- But what batch size to use?
 - Too small = models are too weak
 - Too large = slow to adapt (missing new instances!)

Batch-Incremental Ensemble

Build regular models on batches/windows of instances (typically in a [weighted] ensemble).




- A common approach, surprisingly effective
- Open choice of base classifier (e.g., C4.5, SVM, ...)
- But what batch size to use?
 - Too small = models are too weak
 - Too large = slow to adapt (missing new instances!)
- Can have **sliding** (as opposed to **tumbling**) windows
 - Extreme case: model built every instance
 - But too many batches = too slow

Outline

- 1 Multi-labelled Streaming Data
- 2 k NN for Streams
- 3 Batch-incremental Ensemble
- 4 Incremental base learner**
- 5 Hoeffding Trees
- 6 Neural Networks
- 7 Concept Drift
- 8 Temporal Dependence
- 9 Connections to Sequential Data
- 10 Unlabelled Instances
- 11 Summary

Problem Transformation with Incremental Base Learner

Use an incremental learner (Naive Bayes, SGD, ...) with any problem transformation method (BR, LP, CC, ...)



- Simple deployment, but
- **risk of overfitting** (e.g., with classifier chains),
- concept drift may **invalidate structure**, and a
- **limited choice** of base learner
(must be incremental – e.g., decision tree?)

Outline

- 1 Multi-labelled Streaming Data
- 2 k NN for Streams
- 3 Batch-incremental Ensemble
- 4 Incremental base learner
- 5 Hoeffding Trees**
- 6 Neural Networks
- 7 Concept Drift
- 8 Temporal Dependence
- 9 Connections to Sequential Data
- 10 Unlabelled Instances
- 11 Summary

Rules

For example,

$$\underbrace{X_4 > 1, X_2 \leq 0}_{\text{rule}} \mapsto \underbrace{[1, 0, 1, 0, 0]}_{\mathbf{y}}$$

Can expand rules with the **Hoeffding bound**: for a variable $x \in R$, the Hoeffding bound states that with probability $1 - \delta$, the true mean of X is at least $\bar{x} - \epsilon$, where

$$\epsilon = \sqrt{\frac{R^2 \ln(1/\delta)}{2t}}$$

over t examples.

Rules

For example,

$$\underbrace{X_4 > 1, X_2 \leq 0}_{\text{rule}} \mapsto \underbrace{[1, 0, 1, 0, 0]}_{\mathbf{y}}$$


Can expand rules with the **Hoeffding bound**: for a variable $x \in R$, the Hoeffding bound states that with probability $1 - \delta$, the true mean of X is at least $\bar{x} - \epsilon$, where

$$\epsilon = \sqrt{\frac{R^2 \ln(1/\delta)}{2t}}$$

over t examples. This gives us an indication of when to expand a rule. (You have to select some small value for δ).

Incremental Decision Trees

Recall **decision trees**:



We want to construct a tree **incrementally**, such that the final tree is **with high probability**, identical to that which a traditional (greedy) method would learn.

On a stream $(\mathbf{x}_t, y_t) | t = 1, \dots$, with attributes $\mathbf{X}_\ell = X_1, \dots, X_{D_\ell}$ at leaf ℓ (initially the root):


HOEFFDING TREE INDUCTION ALGORITHM (for each instance):

- ① Sort (\mathbf{x}_t, y_t) into leaf ℓ
- ② Update *sufficient statistics* (to calculate *gain* \bar{G}) in ℓ
- ③ $n_\ell \leftarrow n_\ell + 1$ ▷ *number of examples seen at ℓ*
- ④ If $n_\ell \bmod n_{\min} = 0$:
 - ① Compute $\bar{G}(X_j)$ for all attributes \mathbf{X}_ℓ ▷ *(in leaf ℓ)*
 - ② X_a is the attribute with highest \bar{G} and X_b second highest
 - ③ Compute *Hoeffding bound* $\epsilon = \sqrt{\frac{R^2 \ln(1/\delta)}{2t}}$
 - ④ if $[\bar{G}(X_a) - \bar{G}(X_b)] > \epsilon$, then (knowing that $\Delta G \geq \Delta \bar{G} - \epsilon > 0$ with probability $1 - \delta$):
 - ① turn leaf ℓ into node splitting on X_a
 - ② create leaves ℓ_1, \dots, ℓ_k for all values of $X_a \in \{v_1, \dots, v_k\}$, with initialized sufficient statistics

i.e., with parameters *grace period* n_{\min} , *allowable error* in a split decision δ : 1 - desired probability of choosing correct attribute.

Multi-label Incremental Decision Trees


- Use Hoeffding tree algorithm with **multi-label entropy**
- Examples with *multiple labels* collect at the leaves (we can take the majority labelset, or merge ...)



i.e., if $\{(\mathbf{x}_\ell, [1, 1, 0]), (\mathbf{x}_\ell, [1, 0, 0]), (\mathbf{x}_\ell, [0, 1, 0])\}$ have been seen at leaf ℓ , we could return score $[0.67, 0.67, 0.0]$ (we already have these statistics).


Multi-label Incremental Decision Trees

- A multi-label *incremental* model
- Very *fast* (*single* model for all labels), and usually very *competitive*,



Multi-label Incremental Decision Trees


- A multi-label *incremental* model
- Very *fast* (*single* model for all labels), and usually very *competitive*,



- But Hoeffding bound is **conservative**, tree reluctant to grow, after many examples can remain as root node / stump / majority class classifier, ...
- And when it does grow, it may soon become outdated by **concept drift**

Classifiers at the Leaves

- Place multi-label **classifiers at the leaves** of the tree




where each $\mathbf{h}_\ell : \mathcal{X}^{D^\ell} \rightarrow \mathcal{Y}^{L^\ell}$, i.e., a multi-label classifier dealing with a subset of the input space and a subset of the label space at each leaf ℓ (note: $D^\ell \leq D$, $L^\ell \leq L$).

Hoeffding Adaptive Trees

If the concept changes, do we have to start from scratch? We can use Hoeffding Adaptive Trees (HATs):

- Choose a change-point estimator, e.g., ADWIN, and deploy it at each node
- It will give an alarm when a change at that node is detected
- Chop off the branch from that node, (but we can keep the rest of the tree alive)




(e.g., after adaptation)

Outline


- 1 Multi-labelled Streaming Data
- 2 k NN for Streams
- 3 Batch-incremental Ensemble
- 4 Incremental base learner
- 5 Hoeffding Trees
- 6 Neural Networks**
- 7 Concept Drift
- 8 Temporal Dependence
- 9 Connections to Sequential Data
- 10 Unlabelled Instances
- 11 Summary

Multiple-layer Perceptron



- Direct application, using stochastic gradient descent (SGD) (**incremental**) and varieties, with back-propagation.
- Unlike classifier chains (etc.) we do not need a structure on the output layer – if the inner layer conditionally removes the dependence on the input

Multiple-layer Perceptron




- Direct application, using stochastic gradient descent (SGD) (**incremental**) and varieties, with back-propagation.
- Unlike classifier chains (etc.) we do not need a structure on the output layer – if the inner layer conditionally removes the dependence on the input
- but empirical data-stream comparisons show that SGD/MLPs achieve **poor accuracy** against HTs: they are even more difficult to parameterize for/on a stream!

An Independent Label Space

General approach, decoding inner-layer predictions:

$$\mathbf{y} = f^{-1}(\mathbf{z})$$

$$\mathbf{z} = f(\mathbf{y})$$



- Apply binary relevance, on fewer, independent labels!
- Many methods exist (basis functions, PCA, CCA, factor analysis, deep learning, MLPs, cluster analysis, ...)

No need for back-propagation ...

Removing Dependence

We can just run PCA on the **label space**, make **labels independent**!


$$\mathbf{Z} = \mathbf{W}\mathbf{Y} \quad \triangleright \mathbf{W} \text{ has top } k \text{ components, } k < L$$

$$\mathbf{h} : \mathcal{X} \rightarrow \mathcal{Z} \quad \triangleright \text{model, trained on examples } \{\mathbf{x}_i, \mathbf{z}_i\}_{i=1}^N$$

and then, with input \mathbf{x} ,

$$\mathbf{z} = \mathbf{h}(\mathbf{x}) \quad \triangleright \text{prediction}$$

$$\hat{\mathbf{y}} = \mathbf{W}^{-1}\mathbf{z} \quad \triangleright \text{transform}$$




PCA on the Label Space

- Removes correlation, relatively scalable,
- compresses label space (faster learning for BR, etc.)

but ...

- only **linear correlation**
- doesn't take into account **input space**




source: Wikipedia

- can use **kernels** (kernel PCA ...),
- can use **Canonical Correlation Analysis** (CCA)

Revisiting Meta Labels

Meta labels form an inner layer, e.g., with $K = 2$ meta labels:



where, e.g., $Y_{1,2} \in \{[0, 0], [0, 1], [1, 1]\}$ and $Y_{2,3} \in \{[1, 0], [0, 1]\}$.

- We can achieve $K < L$, and $K < D$,
- without any iterative learning!

Voting Using Meta Labels

Table: An example, for a meta-label on $S_1 = \{1, 2\}$, where $\mathbf{Z}_1 = \{[0, 0], [0, 1], [1, 1]\}$. Posterior distribution $\mathbf{P}_1|\mathbf{x}$ used to obtain $\mathbf{P}'_{S_1} = \mathbf{Z}_1 * \mathbf{P}_1$ (bottom row, left). Using indices in S_1 , we add these values to the final prediction, with the result from other meta-labels ($S_2 = \{2, 3\}$, in this example), to obtain final posterior for labels (bottom row, right).

S_1 :	1	2	\mathbf{P}_1
$[\mathbf{Z}_1]_{:,1}$	0	0	0.0
$[\mathbf{Z}_1]_{:,2}$	0	1	0.9
$[\mathbf{Z}_1]_{:,3}$	1	1	0.1
\mathbf{P}'_{S_1}	0.1	1.0	

j :	1	2	3
\mathbf{P}'_{S_1}	0.1	1.0	
\mathbf{P}'_{S_2}		0.7	0.3
$P(y_j = 1 \mathbf{x})$	0.1	0.85	0.3

Outline


- 1 Multi-labelled Streaming Data
- 2 k NN for Streams
- 3 Batch-incremental Ensemble
- 4 Incremental base learner
- 5 Hoeffding Trees
- 6 Neural Networks
- 7 Concept Drift**
- 8 Temporal Dependence
- 9 Connections to Sequential Data
- 10 Unlabelled Instances
- 11 Summary

Concept Drift: Types

Types of drift:

- 1 Sudden/abrupt
- 2 Incremental
- 3 Gradual

additionally noting the possibility of **reoccurring** drift which may involve any of these types and, noting also the related task of dealing with **outliers**, which is *not* concept drift.



1

Sensors

For example, a sensor may be replaced (sudden); or wear out slowly (incremental), or work only sometimes and increasingly infrequently (gradual).




Concept Drift: Real vs Virtual

Data is drawn from some (unknown) distribution

$$(x_t, y_t) \sim P_t(X, Y) = P_t(Y|X)P_t(X)$$

and either component may drift:




- **Real concept drift**: changes to $P_t(Y|X)$.
- **Virtual drift** changes to $P_t(X)$ only.




Spam Classification

- Near Christmas we receive many more Spam e-mails, but our definition spam has not changed – **virtual drift**.
- We were interested in the latest flight offers, yet now we decide they are unwanted, and mark them as Spam – **real drift**.

Before (left), during (centre) and after (right) concept drift [linear decision boundary]:



Before (left), during (centre) and after (right) concept drift [linear decision boundary]:



- Model becomes invalid as the concept drifts
- **Multi-label** concept drift involves also the *label variables*.

News Articles: Label Activity Over Time




Figure: Multi-label virtual drift: Label frequency / month over time (until about 2007)

Dealing with Concept Drift

Possible approaches to *detecting* and *responding to* concept drift:

- ① Just **ignore it** – batch models must be replaced anyway (we can place higher weight on the most recent models), k NN will eventually have to purge old instances, and SGD adapts constantly (given a learning rate above 0).
- ② Monitor a statistic (e.g., predictive performance, **accuracy** – or a **distribution**) with a **change detector**, and **reset/recalibrate models** when change is detected, e.g., HATs and many ensemble methods.

Dealing with Concept Drift

Possible approaches to *detecting* and *responding to* concept drift:

- 1 Just **ignore it** – batch models must be replaced anyway (we can place higher weight on the most recent models), k NN will eventually have to purge old instances, and SGD adapts constantly (given a learning rate above 0).
- 2 Monitor a statistic (e.g., predictive performance, **accuracy** – or a **distribution**) with a **change detector**, and **reset/recalibrate models** when change is detected, e.g., HATs and many ensemble methods.

In the multi-label case: it is the same, but

- More classifiers are involved
- Models are more complex, e.g., we need $\mathbf{P}_x(Y_1, \dots, Y_L)$ rather than just $\mathbf{P}_x(Y)$.

Detection via Monitoring Accuracy

Drift can be detected via error/accuracy:





Figure: Accuracy through concept drift ($t = 50, \dots, 150$).

Detection via Monitoring Distribution

In the multi-label case, we have a more complex distribution:



i.e., density $p(\epsilon)$ – shape may change over time, and structures may need to be adjusted (the best structure may change)

Multi-label Concept Drift

Consider the **relative frequencies** of labels Y_1 and Y_2 at time t ,

$$\mathbf{C}_t = \frac{1}{t} \mathbf{Y}^\top \mathbf{Y} = \begin{bmatrix} \tilde{p}_1 & \tilde{p}_{1,2} \\ \tilde{p}_{2,1} & \tilde{p}_2 \end{bmatrix}$$

where $\tilde{p}_{1,2} > \tilde{p}_1 \tilde{p}_2$ indicates **marginal dependence**!

Possible **drift** (where $\mathbf{C}_t \neq \mathbf{C}_{t+1}$):

- p_1 increases (label Y_1 relatively **more frequent**)
- p_1 and p_2 both decrease (**label cardinality** decreasing)
- $p_{1,2}$ changes relative to $p_1 p_2$ (change in marginal **dependence** relation between the labels)

Multi-label Concept Drift

And when **conditioned on input** \mathbf{x} , we consider the **relative frequencies**/values of the **errors**, where, e.g., $E_{ij} = (y_j^{(i)} - \hat{y}_j^{(i)})^2$:

$$\mathbf{C}_t = \frac{1}{t} \mathbf{E}^\top \mathbf{E} = \begin{bmatrix} \tilde{p}_1 & \tilde{p}_{1,2} \\ \tilde{p}_{2,1} & \tilde{p}_2 \end{bmatrix}$$

(if **conditional independence**, then $\tilde{p}_{1,2} \approx \tilde{p}_1 \cdot \tilde{p}_2$).

Possible **drift** (where $\mathbf{C}_t \neq \mathbf{C}_{t+1}$):

- p_1 increases (**more errors** on 1-th label)
- p_1 and p_2 both increase (**more errors**)
- $p_{1,2}$ changes relative to p_1, p_2 (change in **conditional dependence** relation)

Outline

- 1 Multi-labelled Streaming Data
- 2 k NN for Streams
- 3 Batch-incremental Ensemble
- 4 Incremental base learner
- 5 Hoeffding Trees
- 6 Neural Networks
- 7 Concept Drift
- 8 Temporal Dependence**
- 9 Connections to Sequential Data
- 10 Unlabelled Instances
- 11 Summary

IID Data Streams

In a stream

$$y_t = h(\mathbf{x}_t) + \epsilon_t$$

it is often assumed that instances arrive in **i.i.d.** form:

- **Identically distributed**,

$$\mathbf{x}_t \sim p(\mathbf{x})$$

for the same p always across all $t = 1, \dots$ within the same concept.

- **Independently distributed**,

$$p(\mathbf{x}_t) = p(\mathbf{x} | \mathbf{x}_{t-1})$$

Is this a valid assumption for data streams?

Measuring Temporal Dependence

We can measure dependence with (for example) the **auto-correlation** function (Pearson's correlation coefficient of a variable with itself, lagged $t + 1$)²,

$$\rho_{Y_t, Y_{t+1}} = \frac{\text{Cov}(Y_t, Y_{t+1})}{\text{Std}(Y_t)\text{Std}(Y_{t+1})} \quad (1)$$

$$= \frac{\sum_{t=1}^{T-1} [(y_t - \bar{y})(y_{t+1} - \bar{y})]}{\sqrt{\sum_{t=1}^{T-1} (y_t - \bar{y})^2 \sum_{t=2}^T (y_t - \bar{y})^2}} \quad (2)$$

$$\approx \frac{\sum_{t=1}^{T-1} [(y_t - \bar{y})(y_{t+1} - \bar{y})]}{\sum_{t=2}^T (y_t - \bar{y})^2} \quad (3)$$

²NB: for large T , the difference in the mean of Y_1, \dots, Y_{T-1} and of Y_2, \dots, Y_T can be ignored, hence Eq. (3).

Measuring Temporal Dependence

We can measure dependence with (for example) the **auto-correlation** function (Pearson's correlation coefficient of a variable with itself, lagged $t + 1$)²,

$$\rho_{Y_t, Y_{t+1}} = \frac{\text{Cov}(Y_t, Y_{t+1})}{\text{Std}(Y_t)\text{Std}(Y_{t+1})} \quad (1)$$

$$= \frac{\sum_{t=1}^{T-1} [(y_t - \bar{y})(y_{t+1} - \bar{y})]}{\sqrt{\sum_{t=1}^{T-1} (y_t - \bar{y})^2 \sum_{t=2}^T (y_t - \bar{y})^2}} \quad (2)$$

$$\approx \frac{\sum_{t=1}^{T-1} [(y_t - \bar{y})(y_{t+1} - \bar{y})]}{\sum_{t=2}^T (y_t - \bar{y})^2} \quad (3)$$

We can generalise to $\rho(k)$ to consider the correlation from y_t and y_{t+k} for any **lag** k (may even be negative).

²NB: for large T , the difference in the mean of Y_1, \dots, Y_{T-1} and of Y_2, \dots, Y_T can be ignored, hence Eq. (3).




Figure: Auto-correlation function on the Electricity dataset, for $k = 1, 2, \dots, 500$; source: Indrė Žliobaitė arXiv:1301.3524v1, Jan 2015.


Outline

- 1 Multi-labelled Streaming Data
- 2 k NN for Streams
- 3 Batch-incremental Ensemble
- 4 Incremental base learner
- 5 Hoeffding Trees
- 6 Neural Networks
- 7 Concept Drift
- 8 Temporal Dependence
- 9 Connections to Sequential Data**
- 10 Unlabelled Instances
- 11 Summary

Naive Bayes

At time t , we see instance x_t , and we wish to make a classification, (e.g., [Naive Bayes](#))

$$\begin{aligned}\hat{y}_t &= \operatorname{argmax}_{y_t \in \{0,1\}} P(y_t, x_t) \\ &= \operatorname{argmax}_{y_t \in \{0,1\}} P(x_t | y_t) P(y_t)\end{aligned}$$




- Not a problem, we maintain an empirical estimate of P s via counting
- At time $t + 1$ we get y_t ; we can now update counts with (x_t, y_t)
- We measure error $\epsilon_t = E(y_t - \hat{y}_t)$, look for drift (e.g., [ADWIN](#)), etc.

Naive Bayes

At time t , we see instance x_t , and we wish to make a classification, (e.g., **Naive Bayes**)

$$\begin{aligned}\hat{y}_t &= \operatorname{argmax}_{y_t \in \{0,1\}} P(y_t, x_t) \\ &= \operatorname{argmax}_{y_t \in \{0,1\}} P(x_t | y_t) P(y_t)\end{aligned}$$




- Not a problem, we maintain an empirical estimate of P s via counting
- At time $t + 1$ we get y_t ; we can now update counts with (x_t, y_t)
- We measure error $\epsilon_t = E(y_t - \hat{y}_t)$, look for drift (e.g., **ADWIN**), etc.

But what if Y_t depends on Y_{t-1} (i.e., **temporal dependence**)?

Hidden Markov Model

If there is temporal dependence, we cannot 'stop' at $P(y_t, x_t)$, because observations \dots, x_{t-1}, x_t are connected via y s:



We write out the joint distribution as³


$$\begin{aligned}\hat{y}_t &= \operatorname{argmax}_{y_t \in \{0,1\}} P(y_t, \mathbf{x}_t) \\ &= \operatorname{argmax}_{y_t \in \{0,1\}} P(y_1) \prod_{\tau=2}^t P(x_\tau | y_\tau) P(y_\tau | y_{\tau-1})\end{aligned}$$

Problem, goes back to y_1 ! Can be a long stream!

³Let $\mathbf{x}_t = [x_1, \dots, x_t]$

Data Stream Classifier with Temporal Dependence

Although with standard **data-stream** assumptions, we are typically dealing with the **filtering** problem:




$$\begin{aligned}\hat{y}_t &= \operatorname{argmax}_{y_t \in \{0,1\}} P(y_t, \mathbf{x}_t, \mathbf{y}_{t-1}) \\ &= \operatorname{argmax}_{y_t \in \{0,1\}} P(x_t | y_t) P(y_t | y_{t-1})\end{aligned}$$

The recursion stops at $t - 1$.

Maximum Entropy Markov Model (MEMM)


A **discriminative** approach:



$$\begin{aligned}\hat{y}_t &= \operatorname{argmax}_{y_t \in \{0,1\}} P(y_t | \mathbf{x}_t) \\ &= \operatorname{argmax}_{y_t \in \{0,1\}} P(y_t | \mathbf{x}_t) P(y_t | y_{t-1}) \\ &= h(\mathbf{x}_t, y_{t-1}) \quad \triangleright \text{Classifier}\end{aligned}$$

Maximum Entropy Markov Model (MEMM)

A **discriminative** approach:




$$\begin{aligned}\hat{y}_t &= \operatorname{argmax}_{y_t \in \{0,1\}} P(y_t | \mathbf{x}_t) \\ &= \operatorname{argmax}_{y_t \in \{0,1\}} P(y_t | \mathbf{x}_t) P(y_t | y_{t-1}) \\ &= h(\mathbf{x}_t, y_{t-1}) \quad \triangleright \text{Classifier}\end{aligned}$$

But what if, say $t = 4$, but we don't have y_{t-2}, y_{t-3} yet?

Forecasting with MEMMs

...




To obtain prediction \hat{y}_t , we can use a **prediction** for \hat{y}_{t-1} :

$$\begin{aligned}\hat{y}_t &= \operatorname{argmax}_{y_t \in \{0,1\}} P(y_t | \mathbf{x}_t) \\ &= \operatorname{argmax}_{y_t \in \{0,1\}} P(y_t | x_t) P(y_t | \hat{y}_{t-1}) \\ &= h(x_t, \hat{y}_{t-1}) \quad \triangleright \text{Classifier}\end{aligned}$$

And so on, back until our last *observed* y , then we propagate forward.

Forecasting with MEMMs

...




To obtain prediction \hat{y}_t , we can use a **prediction** for \hat{y}_{t-1} :

$$\begin{aligned}\hat{y}_t &= \operatorname{argmax}_{y_t \in \{0,1\}} P(y_t | \mathbf{x}_t) \\ &= \operatorname{argmax}_{y_t \in \{0,1\}} P(y_t | \mathbf{x}_t) P(y_t | \hat{y}_{t-1}) \\ &= h(\mathbf{x}_t, \hat{y}_{t-1}) \quad \triangleright \text{Classifier}\end{aligned}$$

And so on, back until our last *observed* y , then we propagate forward. But errors may **propagate** down the *chain*.

(Linear Chain) Conditional Random Fields (CRFs)


To avoid the **error propagation** problem, we may use a CRF:



$$\begin{aligned}\hat{\mathbf{y}} &= \operatorname{argmax}_{\mathbf{y} \in \{0,1\}^T} P(\mathbf{y}|\mathbf{x}) \\ &= \operatorname{argmax}_{\mathbf{y} \in \{0,1\}^T} \prod_{t=2}^T f_t(y_{t-1}, y_t, \mathbf{x})\end{aligned}$$

What is f_t ?

Time Indices \equiv Label Indices



- Labels indices can correspond to steps in time (or space)
- Existing multi-label methodologies can be applied:

$$\hat{\mathbf{y}} = h(\mathbf{x})$$

e.g., binary relevance classifiers, meta labels, classifier chains.

From CRF to Probabilistic Classifier Chains (PCC)

$$\hat{\mathbf{y}}_t = \operatorname{argmax}_{\mathbf{y} \in \{0,1\}^L} p(\mathbf{y}|\mathbf{x}; \mathbf{w}) \quad \triangleright \mathbf{w} \text{ is a set of weights}$$




$$= \operatorname{argmax}_{\mathbf{y} \in \{0,1\}^L} \exp \left\{ \sum_k w_k f_k(y_t, y_{t-1}, \mathbf{x}_t) \right\} \quad \triangleright \text{CRF inference}$$

$$= \operatorname{argmax}_{\mathbf{y} \in \{0,1\}^L} \prod_{t=1}^T \exp \left\{ w_t \cdot f_t(y_t, y_{t-1}, \mathbf{x}_t) \right\} \quad \triangleright e^{a+b} = e^a e^b$$


$$= \operatorname{argmax}_{\mathbf{y} \in \{0,1\}^L} \prod_{t=1}^T f_t(y_t, y_{t-1}, \mathbf{x}_t; \mathbf{w}_t) \quad \triangleright \text{a generic fn } f$$

$$= \operatorname{argmax}_{\mathbf{y} \in \{0,1\}^L} f_1(y_1, \mathbf{x}) \prod_{j=2}^L f_j(y_1, \dots, y_{j-1}, \mathbf{x}) \quad \triangleright \text{PCC!}$$

$$= \operatorname{argmax}_{\mathbf{y} \in \{0,1\}^L} P(y_1|\mathbf{x}) \prod_{j=2}^L P(y_j|y_1, \dots, y_{j-1}, \mathbf{x}) \quad \triangleright \text{where } f_j \approx P_j$$



PCC is a flexible CRF (wrt loss function, base classifier, inference method ...).





Distribution	Type
$p(y, \mathbf{x})$	NB
$p(\mathbf{y}, \mathbf{x})$	HMM
$p(y \mathbf{x})$	LR
$p(\mathbf{y} \mathbf{x})$	CRF
$p(\mathbf{y})$	MRF

and CRFs can be derived as a kind of **classifier chain**.

Forecasting with Classifier Chains

We use classifier-chains type inference for forecasting!



Generate samples $\{\mathbf{y}_t\}_{t=1}^T$,

$$y_1^{(t)} \sim P(y_1|\mathbf{x})$$

$$y_2^{(t)} \sim P(y_2|\mathbf{x}, y_1^{(t)})$$

$$y_3^{(t)} \sim P(y_3|\mathbf{x}, y_1^{(t)}, y_2^{(t)})$$

i.e., $\mathbf{y}_t = [y_1^{(t)}, y_2^{(t)}, y_3^{(t)}]$, i.e., sampling into the future.

Outline

- 1 Multi-labelled Streaming Data
- 2 k NN for Streams
- 3 Batch-incremental Ensemble
- 4 Incremental base learner
- 5 Hoeffding Trees
- 6 Neural Networks
- 7 Concept Drift
- 8 Temporal Dependence
- 9 Connections to Sequential Data
- 10 Unlabelled Instances**
- 11 Summary

Unlabelled Instances in the Stream

What if we *never* get true labels for some \mathbf{x}_t ? In many applications it is unrealistic to expect labels for every instance.

Unlabelled Instances in the Stream

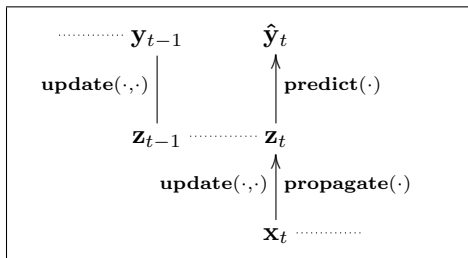
What if we *never* get true labels for some \mathbf{x}_t ? In many applications it is unrealistic to expect labels for every instance. During learning, we can:

- 1 Ignore instances with no label
- 2 Use active learning to get good labels
- 3 Use predicted labels (self-training)
- 4 Use an unsupervised process for example clustering, latent-variable representations.

Semi-Supervised Data Streams

With unsupervised model g to provide a new **representation** of the data:

- 1 $\mathbf{z}_t = g(\mathbf{x}_t)$ ▷ cluster
- 2 $\hat{\mathbf{y}}_t = h(\mathbf{z}_t)$ ▷ predict
- 3 update g with example $(\mathbf{x}_t, \mathbf{z}_t)$
- 4 update h with example $(\mathbf{z}_{t-1}, \mathbf{y}_{t-1})$ (*if \mathbf{y}_{t-1} is available*)




Semi-Supervised Data Streams

With unsupervised model g to provide a new **representation** of the data:

- 1 $\mathbf{z}_t = g(\mathbf{x}_t)$ ▷ cluster
- 2 $\hat{\mathbf{y}}_t = h(\mathbf{z}_t)$ ▷ predict
- 3 update g with example $(\mathbf{x}_t, \mathbf{z}_t)$
- 4 update h with example $(\mathbf{z}_{t-1}, \mathbf{y}_{t-1})$ (if \mathbf{y}_{t-1} is available)

Example



(we can fine tune the inner layer with back propagation).

Outline

- 1 Multi-labelled Streaming Data
- 2 k NN for Streams
- 3 Batch-incremental Ensemble
- 4 Incremental base learner
- 5 Hoeffding Trees
- 6 Neural Networks
- 7 Concept Drift
- 8 Temporal Dependence
- 9 Connections to Sequential Data
- 10 Unlabelled Instances
- 11 Summary

Summary

- Multi-label classification can be adapted to the **data-stream** environment
- This context incurs particular challenges: modelling **label dependence** is important, but this is difficult in a dynamic environment (**concept drift**)
- **Temporal dependence** may exist in data streams: dependence exists across time
- Strong parallels exist between
 - **multi-label learning** (dependence among labels) and
 - **sequence learning** (dependence across time)
- We can adapt existing methods for forecasting and making use of unlabelled instances for training.

Multi-label Streams, Concept Drift, and Sequential Data

Jesse Read



[D&K] IoT Stream Data Mining 2017-2018
December 20, 2017